

Semiparametric robust mean estimations based on the orderliness of quantile averages

Tuban Lee

This manuscript was compiled on May 16, 2023

As one of the most fundamental problems in statistics, robust location estimation has many prominent solutions, such as the trimmed mean, Winsorized mean, Hodges–Lehmann estimator, Huber M -estimator, and median of means. Recent studies suggest that their maximum biases concerning the mean can be quite different, but the underlying mechanisms remain largely unclear. In this article, similar to the mean-median-mode inequality, it is proven that in the context of nearly all common unimodal distributions, there exists an orderliness of symmetric quantile averages with different breakdown points. Further deductions explain why the Winsorized mean and median of means generally have smaller biases compared to the trimmed mean. Building on the U -orderliness, the superiority of the median Hodges–Lehmann mean is discussed.

semiparametric | mean-median-mode inequality | asymptotic | unimodal
| Hodges–Lehmann estimator

In 1823, Gauss (1) proved that for any unimodal distribution with a finite second moment, $|m - \mu| \leq \sqrt{\frac{3}{4}}\omega$, where μ is the population mean, m is the population median, and ω is the root mean square deviation from the mode, M . This pioneering work revealed that despite potential bias with respect to the mean in robust estimates, the deviation remains bounded in unit of a scale parameter under certain assumptions. Bernard, Kazzi, and Vanduffel (2020) (2) further derived asymptotic bias bounds of any quantile for unimodal distributions by reducing this optimization problem to a parametric one, which can be solved analytically. They showed that the population median, m , has the smallest maximum distance to the population mean, μ , among all symmetric quantile averages (SQA_ϵ). Daniell, in 1920, (3) analyzed a class of estimators, linear combinations of order statistics, and identified that ϵ -symmetric trimmed mean (STM_ϵ) belongs to this class. Another popular choice, the ϵ -symmetric Winsorized mean (SWM_ϵ), named after Winsor and introduced by Tukey (4) and Dixon (5) in 1960, is also an L -estimator. Bieniek (2016) derived exact bias upper bounds of the Winsorized mean based on Danielak and Rychlik's work (2003) on the trimmed mean for any distribution with a finite second moment and confirmed that the former is smaller than the latter (6, 7). In 1963, Hodges and Lehmann (8) proposed a class of nonparametric location estimators based on rank tests and, from the Wilcoxon signed-rank statistic (9), deduced the median of pairwise means as a robust location estimator for a symmetric population. Both L -statistics and R -statistics achieve robustness essentially by removing a certain proportion of extreme values. In 1964, Huber (10) generalized maximum likelihood estimation to the minimization of the sum of a specific loss function, which measures the residuals between the data points and the model's parameters. Some L -estimators are also M -estimators, e.g., the sample mean is an M -estimator with a squared error loss function, while the sample median is an M -estimator with an absolute error loss function (10). The

Huber M -estimator is obtained by applying the Huber loss function that combines elements of both squared error and absolute error to achieve robustness against gross errors and high efficiency for contaminated Gaussian distributions (10). Sun, Zhou, and Fan (2020) examined the concentration bounds of Huber M -estimator (11). Mathieu (2022) (12) further derived the concentration bounds of M -estimators and demonstrated that, by selecting the tuning parameter which depends on the variance, Huber M -estimator can also be a sub-Gaussian estimator. The concept of median of means ($MoM_{k,b} = \frac{n}{k}$, k is the number of size in each block, b is the number of blocks) was implicitly introduced several times in Nemirovsky and Yudin (1983) (13), Jerrum, Valiant, and Vazirani (1986), (14) and Alon, Matias and Szegedy (1996) (15)'s works. Given its good performance even for distributions with infinite second moments, MoM has received increasing attention over the past decade (16–23). Devroye, Lerasle, Lugosi, and Oliveira (2016) showed that MoM nears the optimum of sub-Gaussian mean estimation with regards to concentration bounds when the distribution has a heavy tail (21). For a comparison of concentration bounds of trimmed mean, Huber M -estimator, median of means and other relevant estimators, readers are directed to Gobet, Lerasle, and Métivier's paper (2022) (24). Laforgue, Clemencon, and Bertail (2019) proposed the median of randomized means ($MoRM_{k,b}$) (23), wherein, rather than partitioning, an arbitrary number, b , of blocks are built independently from the sample, and showed that MoRM has better non-asymptotic sub-Gaussian property compared to MoM. In fact, asymptotically, the Hodges–Lehmann (H-L) estimator is equivalent to $MoM_{k=2,b=\frac{n}{k}}$ and $MoRM_{k=2,b}$, and they can be seen as the pairwise mean distribution is approximated by the sampling without replacement and bootstrap, respectively. For the asymptotic validity, readers are referred to the foundational works of Efron (1979) (25), Bickel and Freedman (1981, 1984) (26, 27), and Helmers, Janssen, and Veraverbeke (1990)

Significance Statement

In 1964, van Zwet introduced the convex transformation order for comparing the skewness of two distributions. This paradigm shift played a fundamental role in defining robust measures of distributions, from spread to kurtosis. Here, rather than the stochastic ordering between two distributions, the orderliness of quantile averages within a distribution is investigated. By classifying distributions through the signs of derivatives, a series of sophisticated robust mean estimators are deduced. Nearly all common nonparametric robust location estimators are found to be special cases thereof.

T.L. designed research, performed research, analyzed data, and wrote the paper.

The author declares no competing interest.

¹To whom correspondence should be addressed. E-mail: tl@biomathematics.org

72 (28).

Here, the ϵ, b -stratified mean is defined as

$$SM_{\epsilon, b, n} := \frac{b}{n} \left(\sum_{j=1}^{\frac{b-1}{2b\epsilon}} \sum_{i_j=\frac{(2bj-b-1)n\epsilon}{b-1}}^{\frac{(2bj-b+1)n\epsilon}{b-1}} X_{i_j} \right),$$

73 where $X_1 \leq \dots \leq X_n$ denote the order statistics of a sample
 74 of n independent and identically distributed random variables
 75 X_1, \dots, X_n . $b \in \mathbb{N}$, $b \geq 3$. The definition was further refined to
 76 guarantee the continuity of the breakdown point by incorporat-
 77 ing an additional block in the center when $\lfloor \frac{b-1}{2b\epsilon} \rfloor \bmod 2 = 0$,
 78 or by adjusting the central block when $\lfloor \frac{b-1}{2b\epsilon} \rfloor \bmod 2 = 1$ (SI
 79 Text). If the subscript n is omitted, only the asymptotic
 80 behavior is considered. If b is omitted, $b = 3$ is assumed.
 81 $SM_{\epsilon, b=3}$ is equal to STM_{ϵ} , when $\epsilon > \frac{1}{6}$. The basic idea of
 82 the stratified mean, when $\frac{b-1}{2\epsilon} \in \mathbb{N}$, $b \bmod 2 = 1$ is to dis-
 83 tribute the data into $\frac{b-1}{2\epsilon}$ equal-sized non-overlapping blocks
 84 according to their order, then further sequentially group these
 85 blocks into b equal-sized strata and compute the mean of the
 86 middle stratum, which is the median of means of each stratum.
 87 In situations where $i \bmod 1 \neq 0$, a potential solution is to
 88 generate multiple smaller samples that satisfy the equality by
 89 sampling without replacement, and subsequently calculate the
 90 mean of all estimations, the details of determining the sample
 91 size and sampling times are included in the SI Text. Although
 92 the principle is similar to that of the median of means, with-
 93 out the random shift, the result is different from $MoM_{k=\frac{n}{b}, b}$.
 94 Additionally, the stratified mean differs from the mean of the
 95 sample obtained through stratified sampling methods, intro-
 96 duced by Neymean (1934) (29) or ranked set sampling (30),
 97 introduced by McIntyre in 1952, as these sampling methods
 98 are designed to obtain more representative samples or improve
 99 the efficiency of sample estimates, but the sample mean based
 100 on them are not robust. When $b \bmod 2 = 1$, the stratified
 101 mean can be regarded as replacing the other equal-sized strata
 102 with the middle stratum, which, in principle, is analogous to
 103 the Winsorized mean that replaces extreme values with less
 104 extreme percentiles. Furthermore, while the bounds confirm
 105 that the Winsorized mean and median of means outperform
 106 the trimmed mean (6, 7, 21, 24) in worst-case performance,
 107 the complexity of bound analysis makes it difficult to achieve a
 108 complete and intuitive understanding of these results. Also, a
 109 clear explanation for the average performance of them remains
 110 elusive. The aim of this paper is to define a series of semi-
 111 parametric models using the signs of derivatives, reveal their
 112 elegant interrelations and connections to parametric models,
 113 and show that by exploiting these models, a set of sophisti-
 114 cated robust mean estimators can be deduced, which have
 115 strong robustness to departures from assumptions.

116 Quantile average and weighted average

The symmetric trimmed mean, symmetric Winsorized mean,
 and stratified mean are all L -estimators. More specifically,
 they are symmetric weighted averages, which are defined as

$$SWA_{\epsilon, n} := \frac{\sum_{i=1}^{\lceil \frac{n}{2} \rceil} \frac{X_i + X_{n-i+1}}{2} w_i}{\sum_{i=1}^{\frac{n}{2}} w_i},$$

117 where w_i s are the weights applied to the symmetric quantile
 118 averages according to the definition of the corresponding L -
 119 estimators. For example, for the ϵ -symmetric trimmed mean,

$w_i = \begin{cases} 0, & i < n\epsilon \\ 1, & i \geq n\epsilon \end{cases}$, provided that $n\epsilon \in \mathbb{N}$. The mean and
 median are indeed two special cases of the symmetric trimmed
 mean.

To extend the symmetric quantile average to the asymmet-
 ric case, there are two possible definitions for the ϵ, γ -quantile
 average (QA(ϵ, γ, n)), i.e.,

$$\frac{1}{2}(\hat{Q}_n(\gamma\epsilon) + \hat{Q}_n(1 - \epsilon)), \quad [1]$$

and

$$\frac{1}{2}(\hat{Q}_n(\epsilon) + \hat{Q}_n(1 - \gamma\epsilon)), \quad [2]$$

where $\gamma \geq 0$ and $0 \leq \epsilon \leq \frac{1}{1+\gamma}$, $\hat{Q}_n(p)$ is the empirical quantile
 function. For trimming from both sides, [1] and [2] are equiva-
 lent. [1] is assumed in this article unless otherwise specified,
 since many common asymmetric distributions are right skewed,
 and [1] allows trimming only from the right side by setting
 $\gamma = 0$.

Analogously, the weighted average can be defined as

$$WA_{\epsilon, \gamma} := \frac{\int_{\epsilon_0=0}^{\frac{1}{1+\gamma}} QA(\epsilon_0, \gamma) w_{\epsilon_0}}{\int_{\epsilon_0=0}^{\frac{1}{1+\gamma}} w_{\epsilon_0}}.$$

For instance, the ϵ, γ -trimmed mean ($TM_{\epsilon, \gamma}$) is a weighted
 average with a left trim size of $\gamma\epsilon n$ and a right trim size of ϵn ,
 where $w_{\epsilon_0} = \begin{cases} 0, & \epsilon_0 < \epsilon \\ 1, & \epsilon_0 \geq \epsilon \end{cases}$.

Classifying distributions by the signs of derivatives

Let \mathcal{P}_k denote the set of all distributions over \mathbb{R} whose
 moments, from the first to the k th, are all finite. With-
 out loss of generality, all classes discussed in the following
 are subclasses of the nonparametric class of distributions
 $\mathcal{P}_1^k := \{\text{All continuous distribution } P \in \mathcal{P}_k\}$. Besides fully
 and smoothly parameterizing by a Euclidean parameter or
 just assuming regularity conditions, there are many ways to
 classify distributions. In 1956, Stein initiated the problem of
 estimating parameters in the presence of an infinite dimen-
 sional nuisance shape parameter (31). A notable example
 discussed in his groundbreaking work was the estimation of
 the center of symmetry for an unknown symmetric distribution.
 In 1993, Bickel, Klaassen, Ritov, and Wellner published an
 influential semiparametrics textbook (32) and systematically
 classified many common models into three classes: paramet-
 ric, nonparametric, and semiparametric. However, there is
 another old and commonly encountered class of distributions
 that receives little attention in semiparametric literature: the
 unimodal distribution. It is a very unique semiparametric
 model because its definition is based on the signs of deriva-
 tives, i.e., assuming P is continuous, $(f'(x) > 0 \text{ for } x \leq M) \wedge$
 $(f'(x) < 0 \text{ for } x \geq M)$. Let \mathcal{P}_U denote the set of all unimodal
 distributions. Five parametric distributions in \mathcal{P}_U are detailed
 as examples here: Weibull, gamma, Pareto, lognormal and
 generalized Gaussian.

Consider the sign of the derivative of the quantile average
 with respect to the breakdown point, a right-skewed distribu-
 tion is called γ -ordered, if and only if

$$\forall 0 \leq \epsilon \leq \frac{1}{1+\gamma}, \frac{\partial QA_{\epsilon, \gamma}}{\partial \epsilon} \leq 0.$$

The left-skewed case can be obtained by reversing the inequality $\frac{\partial \text{QA}_{\epsilon,\gamma}}{\partial \epsilon} \leq 0$ to $\frac{\partial \text{QA}_{\epsilon,\gamma}}{\partial \epsilon} \geq 0$ and employing the second definition of QA, as given in [2]; for simplicity, it will be omitted in the following discussion. If $\gamma = 1$, the γ -ordered distribution is referred to as ordered.

Furthermore, many common right-skewed distributions are partial bounded, indicating a convex behavior of the QA function when $\epsilon \rightarrow 0$. If assuming convexity further, the second γ -orderliness can be defined as the following for a right-skewed distribution,

$$\forall 0 \leq \epsilon \leq \frac{1}{1+\gamma}, \frac{\partial^2 \text{QA}_{\epsilon,\gamma}}{\partial \epsilon^2} \geq 0 \wedge \frac{\partial \text{QA}_{\epsilon,\gamma}}{\partial \epsilon} \leq 0.$$

Analogously, the ν th γ -orderliness of a right-skewed distribution can be defined as $(-1)^\nu \frac{\partial^\nu \text{QA}_{\epsilon,\gamma}}{\partial \epsilon^\nu} \geq 0 \wedge \dots \wedge -\frac{\partial \text{QA}_{\epsilon,\gamma}}{\partial \epsilon} \geq 0$. If $\gamma = 1$, the ν th γ -orderliness is referred as ν th orderliness. Let \mathcal{P}_O denote the set of all distributions that are ordered and let \mathcal{P}_{O_ν} and $\mathcal{P}_{\gamma O_\nu}$ denote the sets of all distributions that are ν th ordered and ν th γ -ordered, respectively. The following theorems can be used to quickly identify parametric distributions in \mathcal{P}_O , \mathcal{P}_{O_ν} , and $\mathcal{P}_{\gamma O_\nu}$.

Theorem .1. *Any symmetric distribution with a finite second moment is ν th ordered.*

Proof. The assertion follows from the fact that for any symmetric distribution with a finite second moment, all symmetric quantile averages coincide. Therefore, the SQA function is always a horizontal line; the ν th order derivative is zero. \square

As a consequence of Theorem .1 and the fact that generalized Gaussian distribution is symmetric around the median, it is ν th ordered.

Theorem .2. *Any continuous right skewed distribution whose quantile function Q satisfies $Q^{(\nu)}(p) \geq 0 \wedge \dots \wedge Q^{(i)}(p) \geq 0 \wedge Q^{(2)}(p) \geq 0$, $i \bmod 2 = 0$, is ν th γ -ordered, provided that $0 \leq \gamma \leq 1$.*

Proof. Since $(-1)^j \frac{\partial^j \text{QA}_{\epsilon,\gamma}}{\partial \epsilon^j} = \frac{1}{2}((- \gamma)^j Q^j(\gamma\epsilon) + Q^j(1-\epsilon))$, $0 \leq \epsilon \leq \frac{1}{1+\gamma}$, $1 \leq j \leq \nu$, when $j \bmod 2 = 0$, $(-1)^j \frac{\partial^j \text{QA}_{\epsilon,\gamma}}{\partial \epsilon^j} \geq 0$, when $j \bmod 2 = 1$, the strict positivity is uncertain. If assuming $0 \leq \gamma \leq 1$, $(-1)^j \frac{\partial^j \text{QA}_{\epsilon,\gamma}}{\partial \epsilon^j} \geq 0$, since $Q^{(j+1)}(\epsilon) \geq 0$. \square

It is now trivial to prove that the Pareto distribution follows the ν th γ -orderliness, provided that $0 \leq \gamma \leq 1$, since the quantile function of the Pareto distribution is $Q(p) = x_m(1-p)^{-\frac{1}{\alpha}}$, where $x_m > 0$, $\alpha > 0$, so $Q^{(\nu)}(p) \geq 0$ according to the chain rule.

Theorem .3. *A right-skewed continuous distribution with a monotonic decreasing pdf is γ -ordered, if $0 \leq \gamma \leq 1$.*

Proof. A monotonic decreasing pdf means $f'(x) = F^{(2)}(x) \leq 0$. Since $Q'(p) \geq 0$, let $x = Q(F(x))$, then by differentiating both sides of the equation twice, one can obtain $0 = Q^{(2)}(F(x))(F'(x))^2 + Q'(F(x))F^{(2)}(x) \Leftrightarrow Q^{(2)}(F(x)) = -\frac{Q'(F(x))F^{(2)}(x)}{(F'(x))^2} \geq 0$. The desired result is derived from Theorem .2. \square

Theorem .4. *For any random variable X whose probability distribution function belongs to a location-scale family, the distribution is ν th γ -ordered if and only if the family of probability distributions is ν th γ -ordered.*

Proof. Let Q_0 denote the quantile function of the standard distribution without any shifts or scaling, then, after a location-scale transformation, the quantile function is $Q(p) = \lambda Q_0(p) + \mu$, where λ is the scale parameter, μ is the location parameter. According to the definition of the ν th γ -orderliness, the signs of derivatives of the QA function remain the same after this transformation. Since the location-scale transformation can also be performed inversely, the proof is complete. \square

Inequalities related to weighted averages

The bias bound of the ϵ -symmetric trimmed mean is also monotonic for $\mathcal{P}_U \cap \mathcal{P}_2$, as proven in the SI Text using the formulae provided in Bernard et al.'s paper (2). So far, it appears clear that the bias of an estimator is closely related to its degree of robustness. For a right-skewed unimodal distribution, often, $\mu \geq \text{STM}_\epsilon \geq m$. Then analogous to the γ -orderliness, the γ -trimming inequality is defined as $\forall 0 \leq \epsilon_1 \leq \epsilon_2 \leq \frac{1}{1+\gamma}$, $\text{TM}_{\epsilon_1,\gamma} \geq \text{TM}_{\epsilon_2,\gamma}$. A necessary and sufficient condition for the γ -trimming inequality is the monotonic decrease of the bias of the ϵ,γ -trimmed mean as a function of the breakdown point ϵ for a right skewed distribution, proven in the SI Text. γ -orderliness is a sufficient condition for the γ -trimming inequality, as proven in the SI Text, but it is not necessary.

Theorem .5. *For a right-skewed distribution following the γ -trimming inequality, asymptotically, the quantile average is always greater or equal to the corresponding trimmed mean with the same ϵ and γ , provided that $0 \leq \epsilon \leq \frac{1}{1+\gamma}$ and $\gamma \geq 0$.*

Proof. Without loss of generality, assuming the distribution is continuous. According to the definition of the γ -trimming inequality: $\frac{1}{1-\epsilon-\gamma\epsilon+2\delta} \int_{\gamma\epsilon-\delta}^{1-\epsilon+\delta} Q(u) du \geq \frac{1}{1-\epsilon-\gamma\epsilon} \int_{\gamma\epsilon}^{1-\epsilon} Q(u) du$, if $0 \leq \epsilon \leq \frac{1}{1+\gamma}$ and $\gamma \geq 0$, where δ is an infinitesimal positive quantity. Then, rewriting the inequality as $\int_{\gamma\epsilon-\delta}^{1-\epsilon+\delta} Q(u) du - \frac{1-\epsilon-\gamma\epsilon+2\delta}{1-\epsilon-\gamma\epsilon} \int_{\gamma\epsilon}^{1-\epsilon} Q(u) du \geq 0 \Leftrightarrow \int_{1-\epsilon}^{1-\epsilon+\delta} Q(u) du + \int_{\gamma\epsilon-\delta}^{\gamma\epsilon} Q(u) du - \frac{2\delta}{1-\epsilon-\gamma\epsilon} \int_{\gamma\epsilon}^{1-\epsilon} Q(u) du \geq 0$. Since $\delta \rightarrow 0^+$, $\frac{1}{2\delta} \left(\int_{1-\epsilon}^{1-\epsilon+\delta} Q(u) du + \int_{\gamma\epsilon-\delta}^{\gamma\epsilon} Q(u) du \right) = \frac{Q(\gamma\epsilon) + Q(1-\epsilon)}{2} \geq \frac{1}{1-\epsilon-\gamma\epsilon} \int_{\gamma\epsilon}^{1-\epsilon} Q(u) du$, the proof is complete. \square

An analogous result can be obtained in the following theorem.

Theorem .6. *For a right-skewed continuous distribution following the γ -trimming inequality, asymptotically, the Winsorized mean is always greater or equal to the corresponding trimmed mean with the same ϵ and γ , provided that $0 \leq \epsilon \leq \frac{1}{1+\gamma}$ and $0 \leq \gamma \leq 1$.*

Proof. According to Theorem .5, $\frac{Q(\gamma\epsilon) + Q(1-\epsilon)}{2} \geq \frac{1}{1-\epsilon-\gamma\epsilon} \int_{\gamma\epsilon}^{1-\epsilon} Q(u) du \Leftrightarrow \gamma\epsilon(Q(\gamma\epsilon) + Q(1-\epsilon)) \geq \left(\frac{2\gamma\epsilon}{1-\epsilon-\gamma\epsilon} \right) \int_{\gamma\epsilon}^{1-\epsilon} Q(u) du$. Then, if $1 \geq \gamma \geq 0$, $\left(1 - \frac{1}{1-\epsilon-\gamma\epsilon} \right) \int_{\gamma\epsilon}^{1-\epsilon} Q(u) du + \gamma\epsilon(Q(\gamma\epsilon) + Q(1-\epsilon)) \geq$

260 $0 \Rightarrow \int_{\gamma\epsilon}^{1-\epsilon} Q(u) du + \gamma\epsilon Q(\gamma\epsilon) + \epsilon Q(1-\epsilon) \geq \int_{\gamma\epsilon}^{1-\epsilon} Q(u) du +$
 261 $\gamma\epsilon(Q(\gamma\epsilon) + Q(1-\epsilon)) \geq \frac{1}{1-\epsilon-\gamma\epsilon} \int_{\gamma\epsilon}^{1-\epsilon} Q(u) du$, the proof is
 262 complete. \square

263 If assuming γ -orderliness, the result in Theorem .6 can be
 264 extended to the $\gamma > 1$ case, as proven in the SI Text. Replacing
 265 the trimmed mean in the γ -trimming inequality with Win-
 266 sorized mean forms the definition of γ -Winsorization inequality.
 267 γ -orderliness also implies the γ -Winsorization inequality, if
 268 $0 \leq \gamma \leq 1$, as proven in the SI Text.

269 To construct weighted averages based on the γ -orderliness,
 270 let $\mathcal{B}_i = \int_{i\epsilon}^{(i+1)\epsilon} QA(u, \gamma) du$, $ka = k\epsilon + c$, from the
 271 γ -orderliness, $-\frac{\partial QA_{\epsilon, \gamma}}{\partial \epsilon} \geq 0 \Rightarrow \forall 0 \leq a \leq 2a \leq$
 272 $\frac{1}{1+\gamma}, -\frac{(QA(2a, \gamma) - QA(a, \gamma))}{a} \geq 0 \Rightarrow \mathcal{B}_i - \mathcal{B}_{i+1} \geq 0$. Let $\mathcal{B}_i = \mathcal{B}_0$,
 273 then, based on the γ -orderliness, ϵ, γ -block Winsorized mean,
 274 is defined here for comparison in the SI Dataset S1 as

$$BWM_{\epsilon, \gamma, n} := \frac{1}{n} \left(\sum_{i=n\gamma\epsilon+1}^{(1-\epsilon)n} X_i + \sum_{i=n\gamma\epsilon+1}^{2n\gamma\epsilon+1} X_i + \sum_{i=(1-2\epsilon)n}^{(1-\epsilon)n} X_i \right),$$

which is double weighting the leftest and rightest blocks hav-
 ing the size $\gamma\epsilon n$ and ϵn . Since their sizes are different, the
 $0 \leq \gamma \leq 1$ is still necessary for the γ -block Winsorization
 inequality. If γ is omitted, $\gamma = 1$ is assumed. This terminol-
 ogy is the same for other weighted averages. The solutions
 for $i \bmod 1 \neq 0$ are the same as that in SM. From the sec-
 ond γ -orderliness, $\frac{\partial^2 QA_{\epsilon, \gamma}}{\partial^2 \epsilon} \geq 0 \Rightarrow \forall 0 \leq a \leq 2a \leq 3a \leq$
 $\frac{1}{1+\gamma}, \frac{1}{a} \left(\frac{(QA(3a, \gamma) - QA(2a, \gamma))}{a} - \frac{(QA(2a, \gamma) - QA(a, \gamma))}{a} \right) \geq 0 \Rightarrow \mathcal{B}_i -$
 $2\mathcal{B}_{i+1} + \mathcal{B}_{i+2} \geq 0$. So, based on the second orderliness, SM_{ϵ}
 can be seen as assuming $\gamma = 1$, replacing the two blocks,
 $\mathcal{B}_i + \mathcal{B}_{i+2}$ with one block $2\mathcal{B}_{i+1}$. From the ν th γ -orderliness,
 the recurrence relation of the derivatives naturally produces
 the alternating binomial coefficients,

$$\begin{aligned} & (-1)^\nu \frac{\partial^\nu QA_{\epsilon, \gamma}}{\partial \epsilon^\nu} \geq 0 \Rightarrow \forall 0 \leq a \leq \dots \leq (\nu+1)a \leq \frac{1}{1+\gamma} \\ & \frac{(-1)^\nu}{a} \left(\frac{QA(\nu a + a, \gamma)}{a} - \frac{\dots QA(2a, \gamma)}{a} - \frac{QA(\nu a, \gamma)}{a} - \frac{\dots QA(a, \gamma)}{a} \right) \\ & \geq 0 \Leftrightarrow \frac{(-1)^\nu}{a^\nu} \left(\sum_{j=0}^{\nu} (-1)^j \binom{\nu}{j} QA((\nu-j+1)a, \gamma) \right) \geq 0 \\ & \Rightarrow \sum_{j=0}^{\nu} (-1)^j \binom{\nu}{j} \mathcal{B}_{i+j} \geq 0. \end{aligned}$$

Based on the ν th orderliness, the ϵ -binomial mean is introduced
 as

$$BM_{\nu, \epsilon, n} := \frac{1}{n} \left(\sum_{i=1}^{\frac{1}{2}\epsilon^{-1}(\nu+1)-1} \sum_{j=0}^{\nu} \left(1 - (-1)^j \binom{\nu}{j} \right) \mathfrak{B}_{i_j} \right),$$

275 where $\mathfrak{B}_{i_j} = \sum_{l=n\epsilon(j+(i-1)(\nu+1)+1)}^{n\epsilon(j+(i-1)(\nu+1)+1)} (X_l + X_{n-l+1})$. If ν is
 276 not indicated, it is default as $\nu = 3$. Since the alternating
 277 sum of binomial coefficients is zero, when $\nu \ll \epsilon^{-1}$, $\epsilon \rightarrow 0$,
 278 $BM \rightarrow \mu$. If $\frac{1}{2}\epsilon^{-1}(\nu+1) \in \mathbb{N}$, the asymmetry case is
 279 dividing the sample into ϵ^{-1} blocks in the same way as SM and
 280 then further weighting each block using binomial coefficients

($0 \leq \gamma \leq 1$ is needed). The solutions for the continuity of the
 breakdown point and $l \bmod 1 \neq 0$ are the same as that in
 SM and not repeated here. The equality $BM_{\nu=1, \epsilon} = BWM_{\epsilon}$
 holds, and similarly, $BM_{\nu=2, \epsilon} = SM_{\epsilon, b=3}$, when $\gamma = 1$ and
 their respective ϵ s are identical. Interestingly, the biases of the
 $SM_{\epsilon=\frac{1}{9}, b=3}$ and the $WM_{\epsilon=\frac{1}{9}}$ are nearly indistinguishable in
 common asymmetric unimodal distributions such as Weibull,
 gamma, lognormal, and Pareto (SI Text), indicating that their
 robustness to departures from the symmetry assumption is
 practically similar. The reason is that the Winsorized mean is
 using two single quantiles to replace the trimmed parts, not
 two blocks. The following theorems explain this difference.

Theorem .7. For a right-skewed γ -ordered continuous distri-
 bution, asymptotically, the Winsorized mean is always greater
 or equal to the corresponding block Winsorized mean with the
 same ϵ and γ , provided that $0 \leq \gamma \leq 1$.

Proof. From the definitions of BWM and WM, re-
 moving the common part, $\sum_{i=n\gamma\epsilon+1}^{(1-\epsilon)n} X_i$, the state-
 ment requires $\lim_{n \rightarrow \infty} ((n\gamma\epsilon) X_{n\gamma\epsilon+1} + (n\epsilon) X_{n-n\epsilon}) \geq$
 $\lim_{n \rightarrow \infty} \left(\sum_{i=n\gamma\epsilon+1}^{2n\gamma\epsilon} X_i + \sum_{i=n\epsilon}^{2n\epsilon-1} X_{n-i} \right)$. If $0 \leq \gamma \leq 1$, every
 X_i can pair with an X_{n-i+1} to formed a quantile average,
 and the remaining X_{n-i+1} s are all smaller than $X_{n-n\epsilon}$, so
 the inequality is valid. \square

If using single quantiles, based on the second γ -orderliness,
 the stratified quantile mean can be defined as

$$SQM_{\epsilon, \gamma, n} := 4\epsilon \sum_{i=1}^{\frac{1}{4\epsilon}} \frac{1}{2} (\hat{Q}_n((2i-1)\gamma\epsilon) + \hat{Q}_n(1 - (2i-1)\epsilon)),$$

$SQM_{\epsilon=\frac{1}{4}}$ is the Tukey's midhinge (33). In fact, SQM is a
 subcase of SM when $\gamma = 1$ and $b \rightarrow \infty$, so the solution for
 $\frac{1}{\epsilon} \bmod 4 \neq 0$ is the same.

Theorem .8. For a right-skewed second γ -ordered continuous
 distribution, asymptotically, $SQM_{\epsilon, \gamma}$ is always greater or equal
 to the corresponding $BM_{\nu=2, \epsilon, \gamma}$ with the same ϵ and γ , provided
 that $0 \leq \gamma \leq 1$.

Proof. For simplicity, suppose there are $\epsilon^{-1} \in \mathbb{N}$ blocks
 involving in the computation of both SQM and BM. The
 computation of $BM_{\nu=2}$ involves alternating between weight-
 ing and non-weighting, let 0 means the block is assigned
 with a weight of zero, 1 means the block is assigned with a
 weighted of one, the sequence of denoting whether the block
 is weighted or not weighted is: 0, 1, 0, 0, 1, 0, ... Let the se-
 quence be denoted by $a_{BM_{\nu=2}}(j)$, the formula for this sequence
 is $a_{BM_{\nu=2}}(j) = \lfloor \frac{j \bmod 3}{2} \rfloor$. Similarly, the computation of SQM
 can be seen as placing quantiles (p) at the beginning of all
 blocks if $0 < p < \frac{1}{1+\gamma}$, and at the end of all blocks if $p > \frac{1}{1+\gamma}$,
 the sequence of denoting whether the quantile in each block
 is weighted or not weighted is: 0, 1, 0, 1, 0, 1, ... Let the se-
 quence be denoted by $a_{SM}(j)$, the formula for this sequence is
 $a_{SM}(j) = j \bmod 2$. These sequences are also suitable if pairing
 all blocks and quantiles into block average, \mathfrak{B} , and quantile
 average, QA. There are two possible pairing of $a_{BM_{\nu=2}}(j)$ and
 $a_{SM}(j)$, one is $a_{BM_{\nu=2}}(j) = a_{SM}(j) = 1$, another is 0, 1, 0 in
 $a_{BM_{\nu=2}}(j)$ pairing with 1, 0, 1 in $a_{SM}(j)$. By leveraging the
 same principle as Theorem .7 and the second γ -orderliness
 (replacing the two quantile averages with one quantile average
 in the middle), the desired result follows. \square

The biases of $\text{SQM}_{\epsilon=\frac{1}{8}}$, which is based on the second orderliness with a quantile approach, are also very similar to those of $\text{BM}_{\nu=3, \epsilon=\frac{1}{8}}$, which is based on the third orderliness with a block approach, in common asymmetric unimodal distributions (Figure ??).

Hodges–Lehmann inequality and U -orderliness

The Hodges–Lehmann estimator is a very unique robust location estimator due to its definition being substantially dissimilar from conventional symmetric weighted averages. Hodges and Lehmann (8) in their landmark paper *Estimates of location based on rank tests* proposed two methods to compute the H-L estimator, Wilcoxon score R -estimator and median of pairwise means, whose time complexities are $O(n \log(n))$ and $O(n^2)$, respectively. The Wilcoxon score R -estimator is an estimator based on signed-rank test, or R -estimator (8), and was later independently discovered by Sen (34, 35). However, the median of pairwise means is a generalized L -statistic and a trimmed U -statistic (classified by Serfling in his novel conceptualized study in 1984) (36). He also further advanced the understanding by generalizing the H-L kernel, which is $hl_k = \frac{1}{k} \sum_{i=1}^k x_i$, where $k \in \mathbb{N}$ (36). By using the hl_k kernel and the weighted average, it is clear now that the Hodges–Lehmann estimator is also a weighted H-L mean, the definition of which is provided as follows,

$$\text{WeHLM}_{k, \epsilon, \gamma, n} := \text{WA}_{\epsilon_0, \gamma, n} \left((hl_k(X_{N_1}, \dots, X_{N_k}))_{N=1}^{\binom{n}{k}} \right),$$

where $\text{WA}_{\epsilon_0, \gamma, n}(Y)$ denotes the ϵ_0, γ -weighted average with the sequence $(hl_k(X_{N_1}, \dots, X_{N_k}))_{N=1}^{\binom{n}{k}}$ as an input. The asymptotic breakdown point of $\text{WeHLM}_{k, \epsilon, \gamma}$ is $\epsilon = 1 - (1 - \epsilon_0)^{\frac{1}{k}}$ (proven in another relevant paper). The bootstrap method can be used to ensure the continuity of k and therefore the breakdown point. Specifically, let the bootstrap size be denoted by b , then first sampling the original sample $(1 - k + \lfloor k \rfloor)b$ times with the size of each sampling, $\lfloor k \rfloor$, and then subsequently sampling $(1 - \lfloor k \rfloor + k)b$ times with the size of each sampling, $\lfloor k \rfloor$. The corresponding kernels are computed separately, and the pooled sequence is ultimately employed as an input for the WA. The $k = 1$ case is the weighted average. Set the WA in WeHLM as TM_{ϵ_0} , it was named as trimmed H-L mean here (Figure ??, $\epsilon_0 = \frac{15}{64}$). $\text{THLM}_{k=2}$ is close to the Wilcoxon's one-sample statistic investigated by Saleh in 1976 (37), which is first censoring the sample, and then computing the pairwise means. The hl_2 kernel distribution has a probability density function $f_{hl_2}(x) = \int_0^{2x} 2f(t)f(2x-t)dt$ (a result after the transformation of variables), the support of the original distribution is assumed to be $[0, \infty)$ for simplicity. The expected value of the H-L estimator is the positive solution of $\int_0^{\text{H-L}} (f_{hl_2}(s))ds = \frac{1}{2}$. Due to the complexity of this equation, analytically proving the validity of the mean-H-L-median inequality for a distribution is hard. As an example, for the exponential distribution, $f_{hl_2}(x) = 4\lambda^{-2}xe^{-2\lambda^{-1}x}$, $E[\text{H-L}] = \frac{-W_{-1}(-\frac{1}{2e})-1}{2}\lambda \approx 0.839\lambda$, where W_{-1} is a branch of the Lambert W function.

By replacing the H-L kernel with the weighted H-L kernel, which is $whl_k = \frac{1}{k} \sum_{i=1}^k \frac{w_i x_i}{\sum_{i=1}^k w_i}$, the weighted L -statistic can be defined as follows

$$\text{WL}_{k, \epsilon, \gamma, n} := \text{WA}_{\epsilon_0, \gamma, n} \left((whl_k(X_{N_1}, \dots, X_{N_k}))_{N=1}^{\binom{n}{k}} \right).$$

The weighted H-L mean is a special case of the weighted L -statistic when all $w_i = 1$. A complication arises when $w_i \neq 1$; in such cases, regardless of the choice of WA, the weighted L -statistic is not a consistent nonparametric mean estimator when $\epsilon_0 = 0$, so it is not detailed here. If replacing WA in WL with L -estimator, the resulting statistic is referred to as the LL -statistic.

Analogous to the trimming inequality, the Hodges–Lehmann inequality can be defined as $\forall k_2 \geq k_1 \geq 1, m\text{HLM}_{k_2} \geq m\text{HLM}_{k_1}$, where $m\text{HLM}_k$ is setting the WA in WeHLM as median. Since $m\text{HLM}_{k=1} = m$, $m\text{HLM}_{k=2} = \text{H-L}$, $m\text{HLM}_{k=\infty} = \mu$, if a distribution follows the H-L inequality, it also follows the mean-H-L-median inequality. Furthermore, the Hodges–Lehmann inequality is a special case of the γ - U -orderliness, i.e.,

$$\begin{aligned} (\forall k_2 \geq k_1 \geq 1, \text{QHLM}_{k_2, \epsilon, \gamma} \geq \text{QHLM}_{k_1, \epsilon, \gamma}) \vee \\ (\forall k_2 \geq k_1 \geq 1, \text{QHLM}_{k_2, \epsilon, \gamma} \leq \text{QHLM}_{k_1, \epsilon, \gamma}), \end{aligned}$$

where QHLM_k is setting the WA in WeHLM as QA. The direction of the inequality depends on the relative magnitudes of $\text{QA}_{\epsilon, \gamma}$ and μ , since $\text{QHLM}_{k=1, \epsilon, \gamma} = \text{QA}_{\epsilon, \gamma}$ and $\text{QHLM}_{k=\infty, \epsilon, \gamma} = \mu$. U -orderliness is defined as setting $\gamma = 1$.

Theorem .9. U -orderliness implies orderliness.

Proof. The proof is demonstrated by establishing that a distribution exhibiting the U -orderliness property must be ordered. Let $\tilde{\epsilon} = \frac{1}{n}$, when $n \rightarrow \infty$, $\text{SmQHLM}_{k=j, \tilde{\epsilon} \rightarrow 0, n} = \frac{1}{2j} \left(\sum_{i=0}^{j-1} \left(Q\left(\frac{i}{n}\right) + Q\left(\frac{n-i}{n}\right) \right) \right)$, where SmQHLM is setting the WA in WeHLM as SQA. Since Theorem .5 implies that $\mu \leq \text{SQA}_{\epsilon=0}$, when $\tilde{\epsilon} \rightarrow 0$, the statement $\forall k_2 \geq k_1 \geq 1, \text{SmQHLM}_{k_2, \tilde{\epsilon}} \leq \text{SmQHLM}_{k_1, \tilde{\epsilon}}$ is equivalent to the orderliness, i.e., when $n \rightarrow \infty$, $Q\left(\frac{1}{2}\right) \leq \dots \leq \frac{1}{2} \left(Q\left(\frac{i}{n}\right) + Q\left(\frac{n-i}{n}\right) \right) \leq \dots \leq \frac{1}{2} (Q(0) + Q(1))$. \square

Be aware that the U -orderliness itself does not assume any ν th γ -orderliness within the hl_k kernel distribution. The $hl_{k=n-1}$ kernel distribution has n elements, and it can be seen as a location-scale transformation of the original distribution, so it is ν th γ -ordered if and only if the original distribution is ν th γ -ordered according to Theorem .4. If assuming symmetry, the result is trivial since the k -fold convolutions of a symmetric distribution is also symmetric (proved by Laha in 1961) (38). However, proving other cases is challenging. For example, $f'_{hl_2}(x) = 4f(2x)f(0) + \int_0^{2x} 4f(t)f'(2x-t)dt$, the strict negative of $f'_{hl_2}(x)$ is not guaranteed if just assuming $f'(x) < 0$, so, even if the original distribution is monotonic, the hl_2 kernel distribution might be non-monotonic. Also, unlike the pairwise difference distribution, if the original distribution is unimodal, the pairwise mean distribution might be non-unimodal, as demonstrated by a counterexample given by Chung in 1953 and mentioned by Hodges and Lehmann in 1954 (39, 40). If all hl_k kernel distributions, $k \geq 1$, are ν th ordered, then the distribution is ν th U -ordered. From that, the binomial H-L mean (set the WA as BM) can be constructed (Figure ??), while its maximum breakdown point is ≈ 0.065 if $\nu = 3$. A comparison of the biases of $\text{BM}_{\nu=3, \epsilon=\frac{1}{8}}$, $\text{SQM}_{\epsilon=\frac{1}{8}}$, $\text{THLM}_{k=2, \epsilon=\frac{1}{8}}$, $\text{WHLM}_{k=2, \epsilon=\frac{1}{8}}$, $\text{MHLM}_{k=\frac{2 \ln(2) - \ln(3)}{3 \ln(2) - \ln(7)}, \epsilon=\frac{1}{8}}$ (midhinge H-L mean), $m\text{HLM}_{k=\frac{\ln(2)}{3 \ln(2) - \ln(7)}, \epsilon=\frac{1}{8}}$, $\text{THLM}_{k=5, \epsilon=\frac{1}{8}}$,

and $\text{WHLM}_{k=5, \epsilon=\frac{1}{8}}$ is appropriate (Figure ??, SI Dataset S1), given their same breakdown points, with $m\text{HLM}_{k=\frac{\ln(2)}{3\ln(2)-\ln(7)}, \epsilon=\frac{1}{8}}$ exhibiting the smallest biases. This result and Theorem .9 align with Devroye et al. (2016)’s seminal work that MoM is nearly optimal with regards to concentration bounds for heavy-tailed distributions (21), since when k is much smaller than n , the difference between sampling with replacement and without replacement is negligible, $m\text{HLM}_{k,n}$ is asymptotically equivalent to $\text{MoM}_{k,b=\frac{n}{k}}$ if assuming k is a constant. Hence, $\text{MoM}_{k,b=\frac{n}{k}}$ is also based on U -orderliness.

In 1958, Richtmyer introduced the concept of quasi-Monte Carlo simulation that utilizes low-discrepancy sequences, resulting in a significant reduction in computational expenses for large sample simulation (41). Among various numerical sets, Sobol sequences are often favored in quasi-Monte Carlo methods (42). Building upon this principle, in 1991, Do and Hall extended it to bootstrap and found that the quasi-random approach resulted in lower variance compared to other bootstrap Monte Carlo procedures (43). By using a deterministic approach, the variance of $m\text{HLM}_{k,n}$ is much lower than that of $\text{MoM}_{k,b=\frac{n}{k}}$ (SI Dataset S1), when k is small. This highlights the superiority of the median Hodges-Lehmann mean over the median of means, as it not only can provide an accurate estimate for moderate sample sizes, but also allows the use of quasi-bootstrap, where the bootstrap size can be adjusted as needed.

Data Availability. Data for Figure ?? are given in SI Dataset S1. All codes have been deposited in [GitHub](#).

ACKNOWLEDGMENTS. I sincerely acknowledge the insightful comments from the editor which considerably elevated the lucidity and merit of this paper.

20. C Brownlees, E Joly, G Lugosi, Empirical risk minimization for heavy-tailed losses. *The Annals Stat.* **43**, 2507–2536 (2015).
21. L Devroye, M Lerasle, G Lugosi, RI Oliveira, Sub-gaussian mean estimators. *The Annals Stat.* **44**, 2695–2725 (2016).
22. E Joly, G Lugosi, Robust estimation of u -statistics. *Stoch. Process. their Appl.* **126**, 3760–3773 (2016).
23. P Laforgue, S Cléménçon, P Bertail, On medians of (randomized) pairwise means in *International Conference on Machine Learning*. (PMLR), pp. 1272–1281 (2019).
24. E Gobet, M Lerasle, D Métivier, Mean estimation for Randomized Quasi Monte Carlo method. working paper or preprint (2022).
25. B Efron, Bootstrap methods: Another look at the jackknife. *The Annals Stat.* **7**, 1–26 (1979).
26. PJ Bickel, DA Freedman, Some asymptotic theory for the bootstrap. *The annals statistics* **9**, 1196–1217 (1981).
27. PJ Bickel, DA Freedman, Asymptotic normality and the bootstrap in stratified sampling. *The annals statistics* **12**, 470–482 (1984).
28. R Hellmers, P Janssen, N Veraverbeke, *Bootstrapping U-quantiles*. (CWI. Department of Operations Research, Statistics, and System Theory [BS]), (1990).
29. J Neyman, On the two different aspects of the representative method: The method of stratified sampling and the method of purposive selection. *J. Royal Stat. Soc.* **97**, 558–606 (1934).
30. G McIntyre, A method for unbiased selective sampling, using ranked sets. *Aust. journal agricultural research* **3**, 385–390 (1952).
31. C Stein, , et al., Efficient nonparametric testing and estimation in *Proceedings of the third Berkeley symposium on mathematical statistics and probability*. Vol. 1, pp. 187–195 (1956).
32. P Bickel, CA Klaassen, Y Ritov, JA Wellner, *Efficient and adaptive estimation for semiparametric models*. (Springer) Vol. 4, (1993).
33. JW Tukey, , et al., *Exploratory data analysis*. (Reading, MA) Vol. 2, (1977).
34. PK Sen, On the estimation of relative potency in dilution (-direct) assays by distribution-free methods. *Biometrics* pp. 532–552 (1963).
35. M Ghosh, MJ Schell, PK Sen, A conversation with pranab kumar sen. *Stat. Sci.* pp. 548–564 (2008).
36. RJ Serfling, Generalized l -, m -, and r -statistics. *The Annals Stat.* **12**, 76–86 (1984).
37. A Ehsanes Saleh, Hodges-lehmann estimate of the location parameter in censored samples. *Annals Inst. Stat. Math.* **28**, 235–247 (1976).
38. RG Laha, On a class of unimodal distributions. *Proc. Am. Math. Soc.* **12**, 181–184 (1961).
39. J Hodges, E Lehmann, Matching in paired comparisons. *The Annals Math. Stat.* **25**, 787–791 (1954).
40. K Chung, Sur les lois de probabilité unimodales. *COMPTEs RENDUS HEBDOMADAIRES DES SEANCES DE L ACADEMIE DES SCIENCES* **236**, 583–584 (1953).
41. RD Richtmyer, A non-random sampling method, based on congruences, for" monte carlo" problems, (New York Univ., New York. Atomic Energy Commission Computing and Applied . . .), Technical report (1958).
42. IM Sobol', On the distribution of points in a cube and the approximate evaluation of integrals. *Zhurnal Vychislitel'noi Matematiki i Matematicheskoi Fiziki* **7**, 784–802 (1967).
43. KA Do, P Hall, Quasi-random resampling for the bootstrap. *Stat. Comput.* **1**, 13–22 (1991).

1. CF Gauss, *Theoria combinationis observationum erroribus minimis obnoxiae*. (Henricus Dieterich), (1823).
2. C Bernard, R Kazzi, S Vanduffel, Range value-at-risk bounds for unimodal distributions under partial information. *Insur. Math. Econ.* **94**, 9–24 (2020).
3. P Daniell, Observations weighted according to order. *Am. J. Math.* **42**, 222–236 (1920).
4. JW Tukey, A survey of sampling from contaminated distributions in *Contributions to probability and statistics*. (Stanford University Press), pp. 448–485 (1960).
5. WJ Dixon, Simplified Estimation from Censored Normal Samples. *The Annals Math. Stat.* **31**, 385 – 391 (1960).
6. M Bieniek, Comparison of the bias of trimmed and winsorized means. *Commun. Stat. Methods* **45**, 6641–6650 (2016).
7. K Danielak, T Rychlik, Theory & methods: Exact bounds for the bias of trimmed means. *Aust. & New Zealand J. Stat.* **45**, 83–96 (2003).
8. J Hodges Jr, E Lehmann, Estimates of location based on rank tests. *The Annals Math. Stat.* **34**, 598–611 (1963).
9. F Wilcoxon, Individual comparisons by ranking methods. *Biom. Bull.* **1**, 80–83 (1945).
10. PJ Huber, Robust estimation of a location parameter. *Ann. Math. Stat.* **35**, 73–101 (1964).
11. Q Sun, WX Zhou, J Fan, Adaptive huber regression. *J. Am. Stat. Assoc.* **115**, 254–265 (2020).
12. T Mathieu, Concentration study of m -estimators using the influence function. *Electron. J. Stat.* **16**, 3695–3750 (2022).
13. AS Nemirovskij, DB Yudin, *Problem complexity and method efficiency in optimization*. (Wiley-Interscience), (1983).
14. MR Jerrum, LG Valiant, VV Vazirani, Random generation of combinatorial structures from a uniform distribution. *Theor. computer science* **43**, 169–188 (1986).
15. N Alon, Y Matias, M Szegedy, The space complexity of approximating the frequency moments in *Proceedings of the twenty-eighth annual ACM symposium on Theory of computing*. pp. 20–29 (1996).
16. PL Bühlmann, Bagging, subbagging and bragging for improving some prediction algorithms in *Research report/Seminar für Statistik, Eidgenössische Technische Hochschule (ETH)*. (Seminar für Statistik, Eidgenössische Technische Hochschule (ETH), Zürich), Vol. 113, (2003).
17. JY Audibert, O Catoni, Robust linear least squares regression. *The Annals Stat.* **39**, 2766–2794 (2011).
18. D Hsu, S Sabato, Heavy-tailed regression with a generalized median-of-means in *International Conference on Machine Learning*. (PMLR), pp. 37–45 (2014).
19. S Minsker, Geometric median and robust estimation in banach spaces. *Bernoulli* **21**, 2308–2335 (2015).